

SignControl: Multi-Granular Control for Sign Language Video Generation

Xuehan Hou^{1*} Zeyu Zhang^{2*†} Ziye Song^{3*} Huacan Wang⁴ Qizhen Lan⁴
Ronghao Chen⁴ Liuzhou Zhang⁴ Biao Wu⁴ Renda Han⁴ Wenxin Zhang⁴
Hongyang He⁴ Hoi Leong Lee⁴ Guan Huang² Zheng Zhu^{2‡}

¹School of Electronic and Computer Engineering, Peking University, Shenzhen, China

²GigaAI ³NTU ⁴Independent Researcher

*Equal contribution. †Project lead. ‡Corresponding author: zhengzhu@ieee.org.

Abstract

Sign language video generation is a socially significant yet technically demanding task, requiring precise coordination of fine-grained hand movements, facial expressions, and body postures in both spatial and temporal dimensions. While recent diffusion-based video generation models, particularly transformer-based architectures, have advanced the field, existing controllable generation methods face two key challenges in the sign language domain: (1) Control modality granularity mismatch, where heterogeneous control signals such as pose, optical flow, and depth differ in expressive granularity, leading to suboptimal coordination; and (2) Training–inference control inconsistency, where reliance on complete control information during training degrades performance when such information is incomplete at inference. To address these issues, we propose a Hierarchical Multi-modality Control Framework built upon an extended ControlNeXt architecture adapted to the DiT backbone of Wan 2.1 T2V 1.3B. Our method tailors control strategies to each modality’s granularity: fine-grained joint motion control via DWPose, medium-grained pixel-level motion modeling via OnlyFlow with PGMM, and coarse-grained spatial structure control via Depth Anything V2 with Depth LoRA. Furthermore, we introduce a control decay strategy that gradually reduces modality weights during training, enhancing robustness under incomplete control conditions. Experiments on the phoenix-2014.v3 dataset show that our approach substantially outperforms existing methods, including SignGen, in both visual quality and semantic accuracy. Code: <https://github.com/AIGeeksGroup/SignControl>.

1 Introduction

With the flourishing development of deep learning techniques in computer vision, video generation tasks (Ho et al., 2020; Song et al., 2022; Ho et al., 2022; Hou et al., 2026) have evolved

from early simple animation synthesis to generating high-fidelity, semantically consistent complex video content. The emergence of diffusion models has further advanced breakthroughs in this field, particularly models based on Diffusion Transformer architectures (Peebles and Xie, 2023; Mo et al., 2023; Ma et al., 2025) such as Wan 2.1 T2V 1.3B (Wan et al., 2025), which demonstrate powerful capabilities in generating high-quality videos directly from textual descriptions. Among numerous video generation applications, sign language video generation has attracted significant attention due to its unique social value and technical challenges. Unlike general video generation tasks, sign language video generation demands extremely high motion precision and semantic consistency: it is a complex three-dimensional spatial language that encompasses not only hand postures and handshape changes, but also involves multi-layered content including facial expressions, mouth movements, and body postures. These elements must be precisely coordinated in both temporal and spatial dimensions to accurately convey semantic information.

Recent work SignGen (Qi et al., 2025) proposed an end-to-end sign language video generation framework that utilizes latent diffusion models to directly generate sign language videos from text, significantly improving generation quality and naturalness. However, existing controllable video generation methods face two core challenges when applied to the sign language domain. First, the control modality granularity mismatch problem: existing methods such as ControlNet (Zhang et al., 2023a) directly input different types of control signals (pose, optical flow and depth) into multimodal diffusion models without adequately considering the expressive granularity differences across control modalities. Specifically, pose provides discrete joint motions, hand, arm, and finger positions, optical flow provides Dense motion of each pixel, and depth provides 3D structure per pixel. This




Figure 1: **Generated sign language video key frames from text prompts.** The figure displays representative frames generated by our model for different text inputs describing weather-related content. Each row shows 8 key frames extracted from the generated sign language videos, with the corresponding input text prompt displayed below. The results demonstrate the model’s capability to generate diverse sign language sequences for various signers and different textual descriptions, including weather forecasts and news reports.

mismatch of granularity leads to ineffective coordination between different control signals, which affects the quality and accuracy of the generated videos. Second, the training-inference stage control information inconsistency problem: existing methods over-rely on complete control information during training, but in practical inference applications, this control information is usually incomplete or unavailable. This gap between the training and inference stages causes performance degradation during the actual deployment, reducing the practicality of the system.

To address these challenges, we propose a hierarchical multi-modality control-based approach for sign language video generation. Our method builds upon an improved ControlNeXt architecture(Peng et al., 2025), extending it from the original UNet architecture to DiT architecture to accommodate our base model Wan 2.1 T2V 1.3B. To tackle the control modality granularity mismatch problem, we design a hierarchical multi-modality control framework that adopts corresponding processing strategies based on the expressive granular-

ity characteristics of different control modalities: fine-grained control employs DWPose(Yang et al., 2023) to extract pose information and performs precise joint motion control through improved ControlNeXt; medium-grained control employs OnlyFlow(Koroglu et al., 2024) combined with Parsimonious Gaussian Mixture Models (PGMM) input approach to extract optical flow information, capturing pixel-level motion patterns; coarse-grained control employs Depth Anything V2(Yang et al., 2024) to extract depth information and performs spatial structure control through Depth LoRA techniques(Hu et al., 2021; Zhang et al., 2023b). To address the training-inference control information inconsistency problem, we propose a control decay strategy that gradually reduces the weights of various control modalities linearly during training after model convergence stabilization, forcing the model to learn to generate high-quality sign language videos even under incomplete control information conditions.

We conduct comprehensive experiments on the phoenix-2014.v3 dataset, and the results demon-

strate that our method significantly surpasses existing approaches including SignGen in sign language video generation quality and semantic accuracy. The main contributions of this paper can be summarized as follows:

- Propose a hierarchical multi-modality control framework that designs corresponding control strategies based on the granularity characteristics of different control modalities such as pose, depth, and optical flow, achieving more precise and coordinated sign language video generation control.
- Design a control decay strategy that effectively addresses the training-inference stage control information inconsistency problem by gradually reducing control information during training, enhancing model practicality and robustness.
- Successfully extend the ControlNeXt architecture to DiT architecture, enabling it to accommodate modern Transformer-based video generation models and providing a more powerful technical foundation for controllable sign language video generation.

For related work, see Appendix B.

2 Method

In this section, we provide a comprehensive introduction to our method based on hierarchical multi-modal control. We delve into the core architecture of our approach, including the hierarchical multi-modal control framework and control decay strategy, as illustrated in Figure 2.

2.1 Hierarchical Multi-modal Control Signals

To achieve precise control over sign language video generation, we decompose control signals into text conditions and three different granularities of visual control modalities. This hierarchical design can separately handle control requirements at different levels, including semantic content, motion details, temporal coherence, and spatial structure.

Text Condition Text conditions are responsible for specifying the semantic content of sign language videos. We employ a pre-trained T5-XXL text encoder, consistent with the Wan 2.1 T2V model design, to extract text semantic embeddings. This encoder maps input text sequences

to high-dimensional semantic feature space $\mathbf{T} \in \mathbb{R}^{512 \times 4096}$, providing semantic guidance signals for video generation.

Three Visual Control Modalities Targeting the complexity and multi-layered characteristics of sign language videos, we design three different granularities of visual control modalities:

- **Pose Sequences (Fine-grained Control):** Pose information provides joint-level precise motion control and is the most critical control signal in sign language video generation. We employ the DWPose model to extract full-body pose information from input frames, including body joint points, hand keypoints, and facial feature points. DWPose can effectively capture fine gesture changes and facial expressions in sign language expressions, ensuring semantic accuracy of sign language.
- **Optical Flow Sequences (Medium-grained Control):** Optical flow information encodes pixel-level motion patterns, providing temporal coherence guarantees for video generation. We use the OnlyFlow method combined with Parsimonious Gaussian Mixture Models (PGMM) input strategy to extract optical flow information, capturing smooth motion trajectories and dynamic change patterns in sign language actions.
- **Depth Sequences (Coarse-grained Control):** Depth information provides three-dimensional spatial structure constraints of the scene, helping the model understand the spatial hierarchical relationships of sign language actions. We extract depth images through pre-trained Depth Anything V2, which has been specifically optimized for depth estimation tasks and can provide accurate spatial structure information.

Hierarchical Injection and Training Strategy

After obtaining the above control modal images, we input them into corresponding ControlNeXt modules for feature extraction and temporal compression, as shown in Figure 3. The ControlNeXt architecture for each modality is identical, but uses independently trained model weights. Among DiT’s transformer blocks, we adopt a hierarchical injection strategy: pose control features are injected after the 1st block, optical flow control features




Figure 2: **The Overview of SignControl.** The system consists of three training stages: **Stage 1** performs LoRA dataset alignment training where prompts and videos are processed through T5 and VAE encoders respectively, then combined via LoRA adaptation. **Stage 2** conducts ControlNeXt conditional control training, where multi-modal inputs (pose, optical flow, depth) are processed through separate ControlNeXt modules and integrated with text embeddings via CrossNorm operations into DiT blocks with LoRA. **Stage 3** implements LoRA guidance-free adaptation using dropout masking on the original conditions to enable flexible control during inference.

are injected in the 6th block, and depth control features are injected in the 11st block. This hierarchical design enables control information of different granularities to play roles at different depths of the network, achieving comprehensive control from fine-grained gesture details to coarse-grained spatial structure.

The injection process for each control modality adopts statistical alignment mechanisms to ensure feature distribution consistency and controls injection strength through independent learnable scaling factors. During training, we first load pre-trained LoRA weights to achieve coarse alignment, then jointly train each ControlNeXt module, and finally fine-tune LoRA weights to optimize overall performance. This progressive training strategy ensures that the model obtains precise control capabilities while maintaining original generation quality.

2.2 Control Decay Strategy

Unlike existing methods, we adopt a stage-wise training strategy that gradually develops from coarse semantic alignment to precise multi-modal control, ultimately achieving robust generation under incomplete control information.

LoRA Coarse Alignment Stage First, we perform LoRA fine-tuning on the pre-trained DiT model to achieve semantic domain adaptation from general video generation to sign language video generation. At this stage, we only use text conditions for training, with the goal of having the model learn basic patterns and temporal features of sign language actions. LoRA adapters adjust the weights of attention mechanisms and feed-forward networks through low-rank decomposition, enabling the model to understand the basic correspondence between sign language semantics and visual actions. Through this stage of training, the model gains preliminary understanding of sign language video distribution, preparing for the precise control of the next stage.

Multi-modal Control Learning Stage Based on LoRA coarse alignment, we load frozen pre-trained LoRA weights and introduce hierarchical Multi-Modal control modules for joint training. The core of this stage is training each ControlNeXt module so that they can accurately inject different granularities of control signals into specific layers of DiT. We employ cross normalization to ensure distribution consistency between control features and




Figure 3: **Multi-modal input examples.** The figure demonstrates the diverse conditioning inputs used in ControlNeXt training.

latent representations, and dynamically adjust the influence strength of different control modalities through learnable scaling factors. During training, pose, optical flow, and depth control signals are injected at the 1st, 6th, and 11st transformer blocks respectively, forming a control hierarchy from fine-grained to coarse-grained. This hierarchical injection design enables the model to simultaneously handle local gesture details and global spatial structure, achieving comprehensive control over sign language video generation.

Control Decay and Robustness Optimization

To address the problem of control information inconsistency between training and inference stages, we adopt a control decay strategy for fine-tuning LoRA weights in the final stage with frozen ControlNeXt modules. This strategy forces the model to learn to generate high-quality sign language videos under incomplete control information conditions by gradually reducing the availability of control modalities. Specifically, we randomly hide certain frames of some control modalities according to linear decay laws, with their activation probability being:

$$p_m(e) = \max(0.1, 1.0 - \alpha \cdot \max(0, e - e_{stable}))$$

where $m \in \{\text{pose, flow, depth}\}$ represents control modalities, α is the decay coefficient, and

e_{stable} is the number of stable training epochs. Through this progressive decay, the model gradually learns to use text semantics and partially available control information to infer missing motion details, significantly improving robustness in practical applications. Training at this stage enables the model to generate semantically consistent and visually natural sign language videos even when control information is incomplete or inaccurate.

3 Experiments

Implementation Details. The model is trained on the Phoenix-2014.v3 dataset following the three-stage pipeline described in the paper. All experiments are conducted using $4 \times$ NVIDIA H20 GPUs with 96 GB memory per card, using mixed-precision training. We adopt AdamW with a learning rate of 1×10^{-4} for LoRA parameters and 5×10^{-5} for ControlNeXt modules, with weight decay set to 0.01. The global batch size is 64 video clips, each consisting of 21 frames at a latent resolution of $h \times w = 64 \times 64$.

Quantitative Evaluation Results Analysis. Table 1 shows comprehensive quantitative comparison results of our method with various baseline methods on the Phoenix-2014.v3 dataset. Our SignControl achieves significant performance improve-

Table 1: Quantitative comparison results on Phoenix-2014.v3 dataset

Method	BLEU↑	ROUGE-L↑	PSNR↑	SSIM↑	LPIPS↓	FVD↓
VideoLDM	11.2	27.4	12.3	0.59	0.23	1301
VideoPD	16.8	35.6	19.3	0.69	0.31	1132
VideoFusion	14.9	32.1	14.5	0.72	0.18	1032
SignDiff	17.2	38.4	16.2	0.63	0.21	930
SignGen	19.7	45.2	20.2	0.73	0.03	579
SignControl	21.8	47.3	19.8	0.76	0.03	523

Table 2: Ablation on the effectiveness of multi-granular control. Pose provides the strongest semantic guidance, while Flow and Depth further improve temporal coherence and structural consistency.

Pose	Flow	Depth	BLEU↑	ROUGE-L↑	SSIM↑	FVD↓
✓	✗	✗	19.9	45.5	0.741	610
✗	✓	✗	18.7	43.6	0.736	655
✗	✗	✓	18.2	42.9	0.731	690
✓	✓	✗	20.8	46.4	0.752	565
✓	✗	✓	20.4	46.0	0.755	575
✗	✓	✓	19.4	44.7	0.748	605
✓	✓	✓	21.8	47.3	0.760	523

ments on most evaluation metrics. In terms of semantic consistency, our method achieves a BLEU score of 21.8, improving by 10.7% compared to the current state-of-the-art SignGen method’s score of 19.7. The ROUGE-L score improves from 45.2 to 47.3. In terms of video quality, the SSIM metric improves from 0.73 to 0.76, and the LPIPS metric decreases from 0.03 to 0.025, indicating that both structural similarity and perceptual quality of generated videos have improved. Notably, our method is slightly inferior to SignGen on the PSNR metric (19.8 vs 20.2), indicating room for improvement in pixel-level reconstruction accuracy, but overall performance still leads.

Effectiveness of Multi-Granular Control The results in Table 2 demonstrate the complementary roles of different control modalities in sign language video generation. Using a single control signal already yields reasonable performance, with pose clearly providing the strongest semantic guidance, as reflected by higher BLEU and ROUGE-L scores compared to flow or depth alone. Optical flow and depth, while weaker in isolation, contribute important temporal and structural cues, respectively. Combining two controls consistently improves all metrics over single-control settings, indicating that semantic accuracy and visual coherence benefit from multi-source conditioning. The full combination of pose, flow, and depth achieves the best overall performance, reducing FVD substantially while simultaneously improving SSIM

Table 3: Ablation on the effectiveness of the control decay strategy. Smooth schedules (cosine/sigmoid) tend to be slightly more stable than abrupt decay (step), improving robustness under incomplete control at inference.

Decay schedule	BLEU↑	ROUGE-L↑	SSIM↑	LPIPS↓	FVD↓
No decay (always full control)	21.1	46.6	0.754	0.029	548
Linear decay	21.8	47.3	0.760	0.025	523
Cosine decay	22.0	47.5	0.761	0.025	518
Step decay	21.6	47.0	0.757	0.027	532
Sigmoid decay	21.9	47.4	0.761	0.025	520

and language consistency metrics, validating the effectiveness of the proposed multi-granular control design.

Effectiveness of Control Decay Strategy Table 3 analyzes the impact of different control decay schedules on robustness under incomplete control at inference. Without decay, the model tends to overfit to fully available controls, leading to inferior generalization. Introducing gradual decay consistently improves both semantic accuracy and visual quality, with smooth schedules such as cosine and sigmoid decay yielding the most balanced gains across metrics. These strategies outperform abrupt step decay, which slightly degrades temporal stability as reflected by higher FVD and LPIPS. Overall, the results indicate that smooth control decay better regularizes training, enabling the model to rely on partial control signals and text semantics more effectively, thereby enhancing robustness and practical usability.

4 Conclusion

In this work, we introduced SignControl, a hierarchical multi-modality control framework for sign language video generation, addressing the challenges of control modality granularity mismatch and training–inference control inconsistency. By tailoring control strategies to different granularity levels (fine-grained pose, medium-grained optical flow, and coarse-grained depth) and injecting them hierarchically into a DiT-based backbone, our method enables precise and coordinated motion generation. Furthermore, the proposed Control Decay Strategy enhances robustness under incomplete control conditions, improving the practicality of controllable sign language video synthesis. Experimental results on the Phoenix-2014.v3 dataset demonstrate that SignControl achieves state-of-the-art performance in both semantic accuracy and visual quality, outperforming existing methods such as SignGen.

References

- R. Cox and 1 others. 2002. [Tessa: A system to aid communication with deaf people](#). In *Proceedings of the International Conference on Language Resources and Evaluation (LREC)*.
- L. Dong and 1 others. 2024. [Signavatar: Sign language 3d motion reconstruction and generation](#). In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*.
- M. Efthimiou and 1 others. 2012. [Dicta-sign: Sign language recognition, generation and modelling with application in deaf communication](#). In *Proceedings of the International Conference on Language Resources and Evaluation (LREC)*.
- Follow-Your-MultiPose. 2024. Multi-character pose-guided video generation. <https://arxiv.org/abs/2412.16495>.
- S. Gibet and 1 others. 2016. [Sign3d: Building french sign language motion capture corpora for signing avatars](#). In *Proceedings of the International Conference on Language Resources and Evaluation (LREC)*.
- Jonathan Ho, Ajay Jain, and Pieter Abbeel. 2020. [De-noising diffusion probabilistic models](#). *Preprint*, arXiv:2006.11239.
- Jonathan Ho, Tim Salimans, Alexey Gritsenko, William Chan, Mohammad Norouzi, and David J. Fleet. 2022. [Video diffusion models](#). *Preprint*, arXiv:2204.03458.
- Xuehan Hou, Meng Fan, Pengchong Qiao, Zesen Cheng, Yian Zhao, Lei Zhu, Kaiwen Cheng, Chang Liu, and Jie Chen. 2026. [Anchor frame bridging for coherent first-last frame video generation](#). In *The Fourteenth International Conference on Learning Representations*.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2021. [Lora: Low-rank adaptation of large language models](#). *Preprint*, arXiv:2106.09685.
- HumanDiT. 2025. Diffusion transformer for pose-conditioned video generation. <https://arxiv.org/abs/2502.04847>.
- Phillip Isola, Jun-Yan Zhu, Ting Zhou, and Alexei A Efros. 2017. Image-to-image translation with conditional adversarial networks. *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 1125–1134.
- M. Kipp. 2011. [Signgen: End-to-end sign language video generation with latent diffusion](#). *IEEE Transactions on Visualization and Computer Graphics*, 17(4):653–662.
- Mathis Koroglu, Hugo Caselles-Dupré, Guillaume Jeanerret Sanmiguel, and Matthieu Cord. 2024. [On-lyflow: Optical flow based motion conditioning for video diffusion models](#). *Preprint*, arXiv:2411.10501.
- Akide Liu, Zeyu Zhang, Zhixin Li, Xuehai Bai, Yizeng Han, Jiasheng Tang, Yuanjie Xing, Jichao Wu, Mingyang Yang, Weihua Chen, and 1 others. 2025. [Fpsattention: Training-aware fp8 and sparsity co-design for fast video diffusion](#). *arXiv preprint arXiv:2506.04648*.
- Jiabin Luo, Junhui Lin, Zeyu Zhang, Biao Wu, Meng Fang, Ling Chen, and Hao Tang. 2025. [Univid: The open-source unified video model](#). *arXiv preprint arXiv:2509.24200*.
- Xin Ma, Yaohui Wang, Xinyuan Chen, Gengyun Jia, Ziwei Liu, Yuan-Fang Li, Cunjian Chen, and Yu Qiao. 2025. [Latte: Latent diffusion transformer for video generation](#). *Preprint*, arXiv:2401.03048.
- MimicMotion. 2024. Pose-guided high-quality human motion video generation. <https://arxiv.org/abs/2406.19680>.
- Shentong Mo, Enze Xie, Ruihang Chu, Lewei Yao, Lanqing Hong, Matthias Nießner, and Zhenguo Li. 2023. [Dit-3d: Exploring plain diffusion transformers for 3d shape generation](#). *Preprint*, arXiv:2307.01831.
- William Peebles and Saining Xie. 2023. [Scalable diffusion models with transformers](#). *Preprint*, arXiv:2212.09748.
- Bohao Peng, Jian Wang, Yuechen Zhang, Wenbo Li, Ming-Chang Yang, and Jiaya Jia. 2025. [Controlnext: Powerful and efficient control for image and video generation](#). *Preprint*, arXiv:2408.06070.
- PGMM. 2025. Pose-guided motion model for fine-grained sign language video generation. In *ECCV 2024*.
- Fan Qi, Yu Duan, Huaiwen Zhang, and Changsheng Xu. 2025. [Signgen: End-to-end sign language video generation with latent diffusion](#). In *Computer Vision – ECCV 2024*, pages 252–270, Cham. Springer Nature Switzerland.
- Ben Saunders, Necati Cihan Camgoz, and Richard Bowden. 2022. Signing at scale: Learning to co-articulate signs for large-scale photo-realistic sign language production. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 5141–5151.
- S. Saunders and 1 others. 2020. [Progressive transformers for end-to-end sign language production](#). In *Proceedings of the Conference on Computer Vision and Pattern Recognition (CVPR)*.
- Jingwei Shi, Zeyu Zhang, Biao Wu, Yanjie Liang, Meng Fang, Ling Chen, and Yang Zhao. 2025. [Presenta-gent: Multimodal agent for presentation video generation](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 760–773.
- Jiaming Song, Chenlin Meng, and Stefano Ermon. 2022. [Denoising diffusion implicit models](#). *Preprint*, arXiv:2010.02502.

- H. Stoll and 1 others. 2018. [Signsynth: Data-driven sign language video generation](#). In *Proceedings of the International Conference on Computer Vision (ICCV)*.
- The Sign-IDD Team. 2024a. [Sign-idd: Sign language generation using gan](#). *arXiv preprint arXiv:2405.07974*.
- The SignLLM Team. 2024b. [Signllm: Sign languages production large language models](#). *arXiv preprint arXiv:2405.07974*.
- Team Wan, Ang Wang, Baole Ai, Bin Wen, Chaojie Mao, Chen-Wei Xie, Di Chen, Feiwu Yu, Haiming Zhao, Jianxiao Yang, Jianyuan Zeng, Jiayu Wang, Jingfeng Zhang, Jingren Zhou, Jinkai Wang, Jixuan Chen, Kai Zhu, Kang Zhao, Keyu Yan, and 43 others. 2025. [Wan: Open and advanced large-scale video generative models](#). *Preprint*, arXiv:2503.20314.
- Weijie Wang, Jiagang Zhu, Zeyu Zhang, Xiaofeng Wang, Zheng Zhu, Guosheng Zhao, Chaojun Ni, Haoxiao Wang, Guan Huang, Xinze Chen, and 1 others. 2025. Drivegen3d: Boosting feed-forward driving scene generation with efficient video diffusion. *arXiv preprint arXiv:2510.15264*.
- The wSignGen Team. 2024. [wsigngen: Sign language generation using gan](#). In *Proceedings of the ACL Anthology*.
- Lihe Yang, Bingyi Kang, Zilong Huang, Zhen Zhao, Xiaogang Xu, Jiashi Feng, and Hengshuang Zhao. 2024. [Depth anything v2](#). *Preprint*, arXiv:2406.09414.
- Zhendong Yang, Ailing Zeng, Chun Yuan, and Yu Li. 2023. [Effective whole-body pose estimation with two-stages distillation](#). *Preprint*, arXiv:2307.15880.
- Xin Yao, Yunchao Song, and Tao Wang. 2020. Generative temporal models for video synthesis. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 42(8):1981–1993.
- Liuzhou Zhang, Jiarui Ye, Yuanlei Wang, Ming Zhong, Mingju Cao, Wanke Xia, Bowen Zeng, Zeyu Zhang, and Hao Tang. 2025a. Ego3d: Egocentric video generation with long context diffusion. *arXiv preprint arXiv:2512.04515*.
- Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. 2023a. [Adding conditional control to text-to-image diffusion models](#). *Preprint*, arXiv:2302.05543.
- Qingru Zhang, Minshuo Chen, Alexander Bukharin, Nikos Karampatziakis, Pengcheng He, Yu Cheng, Weizhu Chen, and Tuo Zhao. 2023b. [Adalora: Adaptive budget allocation for parameter-efficient fine-tuning](#). *Preprint*, arXiv:2303.10512.
- Zeyu Zhang, Shuning Chang, Yuanyu He, Yizeng Han, Jiasheng Tang, Fan Wang, and Bohan Zhuang. 2025b. Blockvid: Block diffusion for high-quality and consistent minute-long video generation. *arXiv preprint arXiv:2511.22973*.
- Jun-Yan Zhu, Taesung Park, Phillip Isola, and Alexei A Efros. 2017. Unpaired image-to-image translation using cycle-consistent adversarial networks. In *Proceedings of the IEEE international conference on computer vision*, pages 2223–2232.

A Temporal Consistency and Semantic Alignment

Throughout the progressive training process, we pay special attention to maintaining temporal consistency and semantic alignment. Through DiT architecture’s self-attention mechanism, the model can model long-range dependencies between frames, ensuring temporal coherence of generated videos. Meanwhile, text conditions continuously guide the video generation process through cross-attention mechanisms, ensuring accurate correspondence between generated content and input semantics. Our designed stage-wise training strategy not only improves training efficiency but also makes optimization objectives at different stages more explicit, avoiding training instability problems that may arise from simultaneously optimizing multiple complex objectives.

B Related Work

B.1 Sign Language Generation

Sign language generation (SLG) has made significant advancements in facilitating communication between deaf or hard-of-hearing individuals and those who are hearing. Early systems such as Tessa (Cox et al., 2002), DictaSign (Efthimiou et al., 2012), and Sign3D (Gibet et al., 2016) employed 3D models and motion capture technology to enhance the realism of sign language, although this also increased the cost of data acquisition. In recent years, deep learning approaches have become dominant in pose generation for SLG. The Progressive Transformer model (Saunders et al., 2020) was introduced to convert spoken language into 3D sign poses, while transformer-based methods like Text2Gloss2Pose (Saunders et al., 2020) and Gloss2Pose (Saunders et al., 2022) mapped gloss sequences to continuous spatio-temporal pose representations. Additionally, the Text2Gloss2Pose2Video framework (Stoll et al., 2018) extended these models by generating gloss sequences, converting them into skeleton poses, and synthesizing videos. More recent models, such as SignGen (Kipp, 2011), utilize diffusion models to directly generate sign language videos from text, integrating non-manual signals to enhance the naturalness of the output. Recent innovations also include SignLLM (Team, 2024b), a multilingual large language model for generating 3D sign poses, which uses reinforcement learning to accelerate training. The SignAvatar framework (Dong

et al., 2024) leverages variational autoencoders for word-level sign language generation, and models like wSignGen (wSignGen Team, 2024) and SignIDD (Team, 2024a) further improve the realism and accuracy of generated signs, particularly in terms of natural pose generation and video synthesis.

B.2 Pose-Guided Video Generation

Pose-guided video generation has recently gained significant attention in the fields of computer vision and generative models (Zhang et al., 2025a,b; Shi et al., 2025; Wang et al., 2025; Luo et al., 2025; Liu et al., 2025). Early works primarily focused on image-to-image translation models, such as Pix2Pix (Isola et al., 2017) and CycleGAN (Zhu et al., 2017), but they did not address the issue of temporal consistency in video generation. As the complexity of video generation tasks increased, researchers incorporated temporal modeling techniques, such as Long Short-Term Memory (LSTM) networks and Transformer-based models (Yao et al., 2020), to capture temporal dependencies. In the field of pose-guided video generation, recent works have made significant strides. For instance, MimicMotion (MimicMotion, 2024) introduced a pose-guided approach for high-quality human motion video generation, leveraging progressive latent fusion and confidence scaling strategies to improve the naturalness and coherence of generated videos. Follow-Your-MultiPose (Follow-Your-MultiPose, 2024) proposed a novel multi-character video generation framework that uses cross-attention mechanisms and multi-branch control modules, enabling fine-grained control over multiple characters in generated videos. For complex actions and long sequences, HumanDiT (HumanDiT, 2025) proposed a diffusion-based framework for pose-conditioned video generation. The model utilizes prefix latent reference strategies to effectively capture the subtle pose variations and details in long video sequences, improving temporal consistency. Additionally, PGMM (PGMM, 2025) developed a pose-guided motion model that generates fine-grained sign language videos through optical flow warping and pose adaptation modules, demonstrating advantages in generating precise and coherent actions.

C Baseline Method Selection

We select 5 publicly available representative methods for comparative evaluation, which reflect the

latest development trends in video generation. Among them, VideoLDM adopts a latent diffusion model architecture specifically optimized for video generation tasks. VideoPD is based on a probabilistic diffusion framework, showing excellent performance in video quality and temporal consistency. VideoFusion enhances video generation control capability by fusing multiple modal information. SignDiff is a novel method specifically for sign language video generation, adopting ControlNet architecture to achieve controllable generation, demonstrating unique advantages in controlled generation contexts. SignGen, as the current state-of-the-art end-to-end sign language video generation method, uses latent diffusion models to directly generate sign language videos from text, representing the highest level in this field. The selection of these baseline methods fully reflects the technological development trajectory from general video generation to specialized sign language generation, providing comprehensive performance comparison benchmarks for our method.

D Dataset and Preprocessing

We conduct comprehensive experimental evaluation on the Phoenix-2014.v3 dataset, which is one of the most authoritative benchmark datasets in German sign language recognition and generation. The dataset contains approximately 1000 sentences of high-quality sign language videos, mainly covering common sign language expressions in weather forecast scenarios, with each video annotated with German text and corresponding sign language vocabulary sequences. For video data preprocessing, we adopt a standardized processing pipeline to ensure training consistency and stability. Original videos are first resized to the model-recommended resolution of 480×832, which maintains detailed information of sign language gestures while considering model generation quality. The model uses a fixed length of 81 frames as input sequence, projecting videos to latent space through a pre-trained VAE encoder, obtaining latent representations with dimensions [21, 16, 60, 104]. This design achieves an 8× downsampling ratio, significantly reducing computational complexity while preserving important spatial and temporal information.

E Future Work

For future work, several directions are worth exploring. First, extending the framework to multi-

signer and cross-lingual sign language generation could broaden its applicability, enabling adaptation to different signing styles, regional variations, and multiple sign languages while preserving semantic consistency. Second, explicitly incorporating non-manual signals (such as facial expressions, mouth shapes, and gaze direction) with precise temporal alignment may enhance semantic fidelity, given their critical grammatical and emotional roles. Third, designing adaptive control selection mechanisms that dynamically adjust modality reliance based on input quality could improve robustness in real-world scenarios where certain control signals are noisy. Finally, leveraging larger-scale datasets and exploring multi-task learning with related tasks such as sign language recognition or translation could further boost the model’s generalization and transferability, while continual learning strategies may allow long-term adaptation to evolving linguistic patterns without full retraining.